

LLaMA 3.2 Project Snapshot Report

Scope:

- Benchmark comparison (HellaSwag, ARC, GSM8K)
- Efficiency comparison (TTFT, latency/token, throughput, RAM)


Data Sources (latest matching runs):

- LLaMA-3.2-1B benchmark: benchmark_benchmark_20260415T005401Z_37871af3.json
- LLaMA-3.2-1B efficiency: efficiency_efficiency_20260415T031459Z_91af67a8.json
- LLaMA-3.2-3B benchmark: benchmark_benchmark_20260415T025445Z_d35c46e2.json
- LLaMA-3.2-3B efficiency: efficiency_efficiency_20260415T031329Z_19f7f4ef.json
- Phi-3-mini benchmark: benchmark_benchmark_20260415T014251Z_f17dafa1.json
- Phi-3-mini efficiency: efficiency_efficiency_20260415T031545Z_7376440c.json


Generated at: 2026-04-15T03:18:40.932126+00:00

Benchmark Accuracy and Runtime (limit=20, fp16, mps)


HellaSwag Accuracy




ARC Accuracy



GSM8K Accuracy




Total Benchmark Duration (sec)




Efficiency Metrics (num_prompts=5, timed_runs=2, max_new_tokens=64, fp16, mps)


TTFT Mean (ms)




Latency per Token Mean (ms)



Throughput Mean (tok/s)



RAM After Run (GB)



Interpretation Snapshot

Overall Accuracy Ranking:

1. Phi-3-mini (0.550)
2. LLaMA-3.2-3B (0.483)
3. LLaMA-3.2-1B (0.250)

Benchmark Runtime Ranking (faster is better):

1. LLaMA-3.2-1B (108.1s)
2. LLaMA-3.2-3B (241.3s)
3. Phi-3-mini (302.7s)

Efficiency Highlights:

- Fastest TTFT: LLaMA-3.2-1B
- Best throughput: LLaMA-3.2-1B
- Lowest latency/token: LLaMA-3.2-1B

Caveat:

- These are small-sample benchmark runs (limit=20). Use larger limits for final claims.